


E-Commerce Data Processing Pipeline

Medallion Architecture ETL Pipeline on AWS



1. Executive Summary

This project implements an end-to-end data engineering pipeline that ingests, cleans, transforms, and analyzes e-commerce transactional data at scale. The pipeline follows a Medallion architecture (Bronze → Silver → Gold) built on AWS, using PySpark for distributed processing, AWS Glue for managed ETL orchestration, S3 as the data lake, and Power BI for business-facing dashboards. The goal is to convert raw, messy order/customer/product data into governed, analytics-ready datasets that support revenue reporting, customer behavior analysis, and operational monitoring.

Key outcomes:

- Automated ingestion of 4 core entities (orders, customers, products, clickstream events) into a partitioned S3 data lake.
- PySpark transformation jobs for deduplication, null handling, type enforcement, and referential integrity checks.
- A star-schema Gold layer optimized for BI consumption via Athena / Redshift Serverless.
- An interactive Power BI dashboard surfacing revenue trends, category performance, and regional order distribution.
- Pipeline orchestration and scheduling via Apache Airflow with retry and alerting logic.

2. Problem Statement & Objectives

E-commerce platforms generate high-velocity, high-volume data across multiple disconnected systems — order management, customer profiles, product catalogs, and web clickstream logs. Without a unified pipeline, this data remains siloed, inconsistent, and unusable for timely business decisions. This project addresses that gap by building a scalable, fault-tolerant pipeline that standardizes and aggregates this data for analytics.

Objectives:

- Design a layered data lake (Bronze/Silver/Gold) that separates raw, cleaned, and business-ready data.
- Build reusable, parameterized PySpark ETL jobs deployable as AWS Glue jobs.
- Enforce data quality rules: schema validation, deduplication, null-handling, and referential checks.
- Model a star schema in the Gold layer to support fast, low-cost SQL analytics via Athena.
- Deliver a self-service Power BI dashboard for revenue, product, and regional performance tracking.

3. Architecture Overview

The pipeline follows a Medallion architecture, a layered data lake design pattern that incrementally improves data quality and structure at each stage:

Layer	Purpose	Format / Store
Bronze	Raw, immutable copy of source data, ingested as-is	Parquet on S3 (partitioned by date)
Silver	Cleaned, deduplicated, type-cast, conformed data	Parquet on S3 (partitioned by entity)
Gold	Aggregated, business-level star schema for BI/SQL	Parquet on S3 + Redshift Serverless / Athena tables

E-Commerce Data Processing Pipeline — Architecture

Medallion Architecture (Bronze → Silver → Gold) on AWS




Figure 1: End-to-end pipeline architecture, from source systems to Power BI.

Orchestration

Apache Airflow schedules and coordinates the daily pipeline run: it triggers the S3 landing job, launches the Bronze→Silver and Silver→Gold AWS Glue jobs in sequence, runs data quality checks between stages, and sends failure alerts. Each task is idempotent and can be safely re-run for backfills.

4. Data Sources & Schema

Four core entities are ingested daily as CSV/JSON drops into the S3 landing zone, then converted to Parquet in the Bronze layer:

Entity	Grain	Key Fields
orders	1 row per order line	order_id, customer_id, product_id, order_date, qty, unit_price, status
customers	1 row per customer	customer_id, name, email, signup_date, region, segment
products	1 row per SKU	product_id, category, sub_category, brand, cost_price, list_price
clickstream	1 row per event	session_id, customer_id, event_type, product_id, event_ts

The Gold layer models this into a star schema: **fact_orders** at the center, with **dim_customer**, **dim_product**, and **dim_date** as conformed dimensions — optimized for fast aggregation queries in Athena / Redshift Serverless.

5. Technology Stack

Layer	Technology
Storage	AWS S3 (data lake, partitioned Parquet)
Processing	PySpark (Apache Spark), AWS Glue (managed Spark jobs)
Orchestration	Apache Airflow
Query / Warehouse	AWS Athena, Redshift Serverless
Language	Python, SQL
BI / Reporting	Power BI
Version Control	Git / GitHub

6. ETL Pipeline Implementation

6.1 Bronze Layer — Raw Ingestion (PySpark on AWS Glue)

Raw CSV files land in an S3 landing bucket. A Glue job reads them with an enforced schema and writes immutable, date-partitioned Parquet to the Bronze bucket — no transformation logic, only type casting and ingestion metadata.

```
from pyspark.sql import SparkSession
from pyspark.sql import functions as F

spark = SparkSession.builder.appName("bronze_ingest").getOrCreate()

RAW_PATH = "s3://ecom-landing/orders/"
BRONZE_PATH = "s3://ecom-datalake/bronze/orders/"

schema_orders = "order_id STRING, customer_id STRING, product_id STRING, " \
                 "order_date STRING, qty INT, unit_price DOUBLE, status STRING"

df = (spark.read
      .option("header", True)
      .schema(schema_orders)
      .csv(RAW_PATH))

df = (df
      .withColumn("ingest_ts", F.current_timestamp())
      .withColumn("ingest_date", F.to_date("order_date")))

(df.write
 .mode("append")
 .partitionBy("ingest_date")
 .parquet(BRONZE_PATH))
```

6.2 Silver Layer — Cleaning & Conforming

The Silver job deduplicates records, drops nulls in required fields, standardizes types and casing, and enforces referential integrity against dimension tables before writing cleaned Parquet.

```
bronze = spark.read.parquet("s3://ecom-datalake/bronze/orders/")

silver = (bronze
 .dropDuplicates(["order_id"])
 .filter(F.col("customer_id").isNotNull() & F.col("product_id").isNotNull())
 .withColumn("status", F.lower(F.trim(F.col("status"))))
 .withColumn("order_amount", F.round(F.col("qty") * F.col("unit_price"), 2))
 .filter(F.col("qty") > 0)
 )

valid_customers =
spark.read.parquet("s3://ecom-datalake/silver/customers/").select("customer_id")
silver = silver.join(valid_customers, on="customer_id", how="inner")

(silver.write
 .mode("overwrite")
 .partitionBy("ingest_date")
 .parquet("s3://ecom-datalake/silver/orders/"))
```

6.3 Gold Layer — Star Schema Aggregation

The Gold job builds business-ready fact and dimension tables, pre-aggregating common metrics (daily revenue, category totals) to keep BI queries fast and cheap.

```
fact_orders = (silver
    .groupBy("order_date", "customer_id", "product_id")
    .agg(F.sum("order_amount").alias("revenue"),
        F.sum("qty").alias("units_sold"),
        F.count("order_id").alias("order_count"))
)

daily_category_revenue = (silver
    .join(dim_product, "product_id")
    .groupBy("order_date", "category")
    .agg(F.sum("order_amount").alias("revenue"))
)

(fact_orders.write.mode("overwrite")
    .partitionBy("order_date")
    .parquet("s3://ecom-datalake/gold/fact_orders/"))
```

7. SQL Analytics Layer

The Gold layer is registered as external tables in AWS Glue Data Catalog and queried via Athena. Sample analytical queries used to power the Power BI dashboard:

Top-performing categories by revenue:

```
SELECT
    p.category,
    SUM(f.revenue) AS total_revenue,
    SUM(f.units_sold) AS total_units,
    COUNT(DISTINCT f.customer_id) AS unique_customers
FROM gold.fact_orders f
JOIN gold.dim_product p ON f.product_id = p.product_id
WHERE f.order_date >= date_add('month', -3, current_date)
GROUP BY p.category
ORDER BY total_revenue DESC;
```

Month-over-month revenue growth:

```
WITH monthly AS (
    SELECT date_trunc('month', order_date) AS month,
           SUM(revenue) AS revenue
    FROM gold.fact_orders
    GROUP BY 1
)
SELECT
    month,
    revenue,
    revenue - LAG(revenue) OVER (ORDER BY month) AS mom_change,
    ROUND(100.0 * (revenue - LAG(revenue) OVER (ORDER BY month))
          / LAG(revenue) OVER (ORDER BY month), 2) AS mom_pct
FROM monthly
ORDER BY month;
```

8. Power BI Dashboard

The Gold layer feeds a Power BI dashboard refreshed daily via Athena connector, giving stakeholders a single view of revenue trends, category performance, regional demand, and payment behavior.




Figure 2: E-Commerce Sales & Operations dashboard built on the Gold layer.

9. Results & Business Impact

Metric	Before Pipeline	After Pipeline
Time to produce revenue report	~2 days (manual Excel)	< 15 minutes (automated)
Data quality issues (dup/null rows)	~7% of records	< 0.5% of records
Analyst query latency	Full table scan on raw CSV	Sub-5s on partitioned Parquet via Athena
Reporting cadence	Weekly, manual	Daily, automated refresh

10. Challenges & Learnings

- Schema drift in raw CSV drops required enforcing an explicit schema at Bronze ingestion rather than relying on inference.
- Late-arriving customer records caused referential join failures in Silver — resolved with a nightly reconciliation job.
- Small-file problem on S3 from frequent writes was mitigated by compacting partitions during the Gold job.
- Athena query costs were reduced significantly by partitioning on order_date and using columnar Parquet instead of CSV.

11. Conclusion

This pipeline demonstrates a production-style approach to e-commerce data engineering: a governed, layered data lake; reusable PySpark/Glue transformation logic; a query-optimized star schema; and a business-facing BI layer. The design is cloud-native, cost-efficient, and extensible — new data sources (e.g. marketing spend, support tickets) can be added as additional Bronze ingestion jobs without disrupting downstream Silver/Gold logic.